
Hot Topics

Software-defined Networking

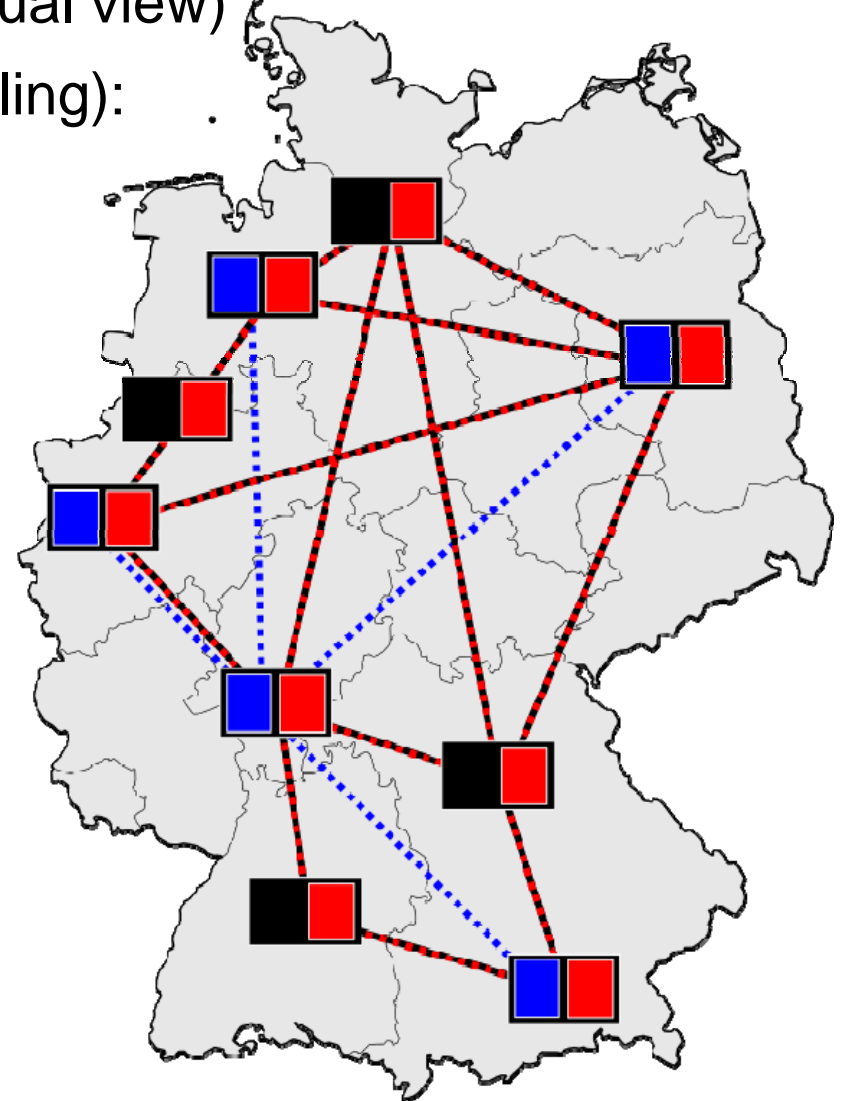
Contents - Hot Topics - SDN

- History
- Control and Data Plane Separation
- Network Virtualization
- Overview of Control Plane
- Network Functions Virtualization

History

Network Virtualization

- representation of one or more logical network topologies on the same infrastructure (network slicing with individual view)
- many different instantiations (using tunneling):
 - Virtual LAN (VLAN)
 - Stacked VLANs (QinQ)
 - Virtual Extensible LAN (VXLAN)
 - NVGRE (Network Virtualization using GRE)

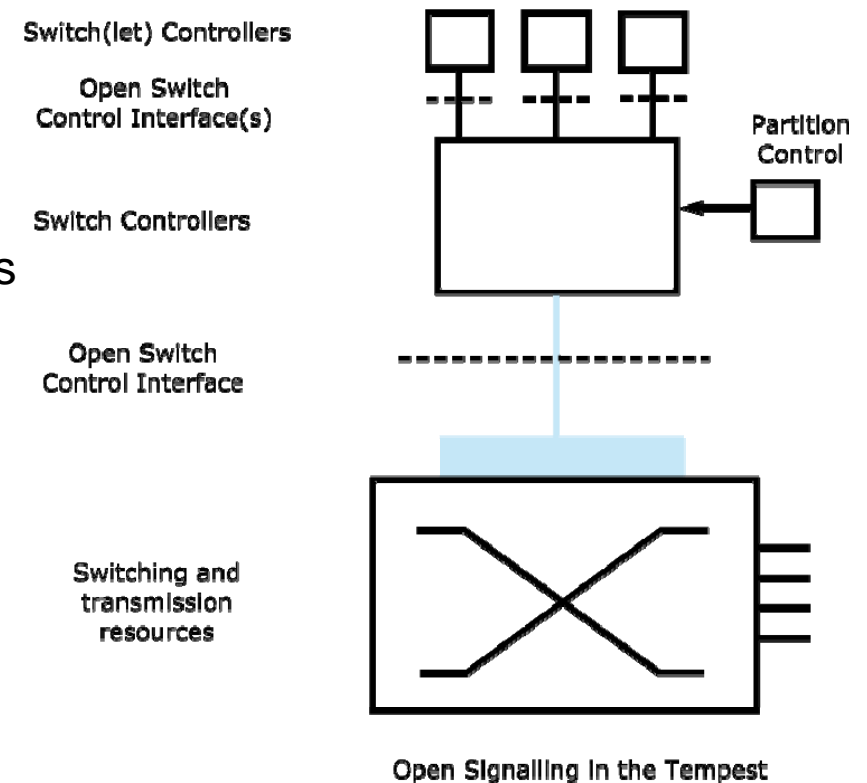


Network Virtualization - Benefits

- Sharing:
 - instantiate multiple logical routers on a single platform
 - requires resource isolation in CPU, memory, bandwidth, forwarding tables, etc.
- Customizability:
 - customizable routing and forwarding software
 - general-purpose CPUs for control plane
 - network processors and FPGAs/ASICs for data plane

Network Virtualization - Examples

- Tempest: Switchlets (1998)
 - virtualization of switches
 - separation of control framework from switches
 - similar to OpenFlow
 - problem: requires standardization, adoption and deployment of new hardware

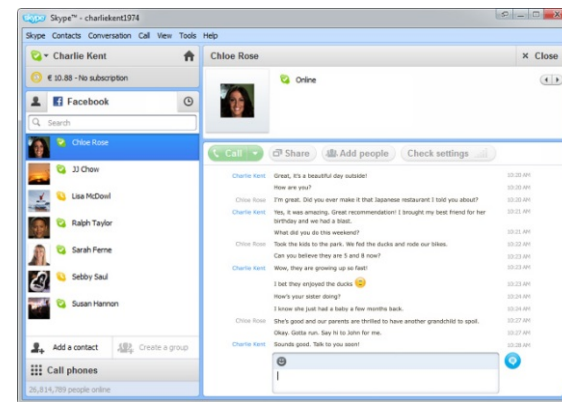



- VINI: Virtual Network Infrastructure (2006)
 - virtualization of the network infrastructure
 - separation of data and control planes
 - bridging the gap between small-scale experiments and live deployment at scale



Network Virtualization - Examples

- Cabo: Concurrent Architectures are Better than One (2007)
 - virtualization of services
 - infrastructure providers (maintaining routers, links, datacenters) can operate independently from service providers (offering end-to-end/over-the-top services, e.g. VPN, Videotelephony)



Legacy of Virtual Networks for SDN:

- separating services from infrastructure (later: NFV)
- logic network topologies on top of physical infrastructure

Plane Separation - Examples



1) Separate control channel: IETF FORCES (RFC5810, 2003)

- forwarding elements (FE) can be controlled by multiple control elements (CE) via a standard control channel (FORCE):

- → forwarding packets, metering, shaping, NATing, traffic classification
- Problem: requires standardization, adoption and deployment of new hardware

2) In-band signaling: Routing Control Platform (2004)

- reuse existing protocols as control channels
- one RCP per AS compute routes on behalf of the routers
- reuse existing routing protocol (BGP) to signal those routes



Plane Separation - Examples

2) In-band signaling: Routing Control Platform (2004)

- to routers it looks like a normal route announcement from another router (in fact to RCP)
- deployment easier using this approach → no standardization, adoption, deployment necessary
- Problem: control constrained by what existing protocol (i.e. BGP) supports


3) Customized hardware: Ethane (2007)

- allows direct enforcement of a single, fine-grained network policy
- domain controller computes flow table entries (to be installed in switches) based on access control policies
- Problem: requires custom switches supporting Ethane protocol: OpenWrt, NetFPGA, Linux

Plane Separation - Examples

4) Hardware with open interfaces: OpenFlow (2008)

- best of both worlds: operate on existing protocols without customizing hardware → OpenFlow
- taking existing capabilities of hardware switches (flow tables already implemented)
- switch exposes flow table through simple OpenFlow protocol
- vendors can keep platform closed, but expose an open interface to control forwarding table
- Switch matches subset of packet header fields:
 - Switch port
 - MAC src/dst
 - Ethernet type
 - VLAN ID/priority
 - IP src/dst + IP protocol + IP ToS bits
 - TCP/UDP sport/dport



Control and Data Plane Separation

Control and Data Plane Separation

1) control plane:

- logic for controlling forward behavior
- usually implemented in general-purpose hardware for easy modification/adaptation
- "brain" of the network
- examples: routing protocols, network middlebox configuration

2) data plane:

- forward traffic according to control plane logic
- often implemented in special hardware for increased throughput: ASICs, FPGAs
- examples: IP forwarding, Layer 2 switching

Plane Separation - Reasoning


Why separation?

- greater flexibility: new services introduced more easily
- independent evolution and development: software can evolve independently of the hardware
- accelerating innovation in existing networks:
 - control logic is not tied to hardware → can be patched in software
 - introducing technologies more rapidly without consensus standardization
 - user driven innovation
- control from high-level software programs: high level of abstraction, debug/check more easily
- network-wide view of controller: easier to infer (and reason) about network behaviour

Plane Separation - Opportunities

1) mitigating DoS (Denial-of-service) attacks by filtering attack traffic:


- measurement system detects attack and identifies entry point and victim of attack
- control plane installs a null route at the entry point
- → data plane drops offending traffic at the entry point
- e.g. AT&T IRSCP (commercial RCP)



Plane Separation - Opportunities

2) Interdomain routing: Constrained policies

- artificially constrained routes of today's interdomain routing protocol BGP
 - route selection based on a fixed set of steps
 - limited knobs to control in-/outbound traffic
 - incorporating additional information (reputation of route, time of day) problematic
- → route selection on a richer set of policies: route controller can directly update state in forwarding elements




2.1) Maintenance dry-out

- e.g. planned maintenance on an edge router: directly tell ingress router to use egress2 instead of egress1
- much more difficult using traditional routing to adjust route of a single router: tuning route weights (indirect way)

Plane Separation - Opportunities

2.2) Egress selection


- customer-controlled egress selection if multiple paths to reach same destination (depending on source rather than destination)
- possibly even giving customers control over the decision
- today: routing traffic based on destination prefix



Plane Separation - Opportunities

2.2) Enhanced security

- anomaly detection: detecting suspicious/bogus routes
- prefer "familiar" routes over unfamiliar ones using reputation of routes
- today: no easy way to incorporate route-reputation into route selection



Plane Separation - Opportunities

3) Datacenters

- cost: expensive vendor switched vs. inexpensive commodity/off-the-shelf switches
- flexible control: tailor network for services, quickly improve and innovate

Other opportunities:

- dynamic access control
- seamless mobility/migration → decreased latency due to VM migration
- server load balancing
- network virtualization
- using multiple wireless access points
- energy-efficient networking
- adaptive traffic monitoring
- DoS attack detection/mitigation (see example 1)

Plane Separation - Challenges

1) Scalability

- one control element responsible for many (thousands) forwarding elements
- RCP: storing route and computing route decisions for every (potentially thousands) router
- → single point of failure
- Solution: aggregation, hierarchy, redundancy, distributed coordination

2) Reliability/Security: What happens when a controller fails or is compromised?

- multiple identical controllers ("hot spare"), connected to same nodes, performing same route calculation → implicit consistency
- takes over if primary fails

Plane Separation - Challenges

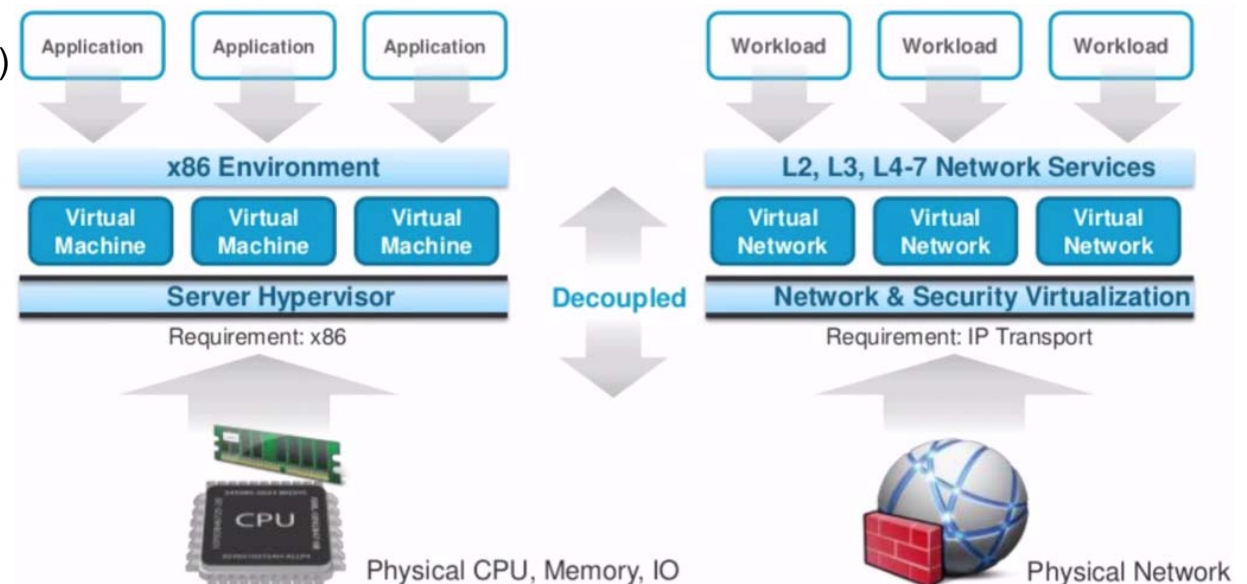
3) Consistency

- implicit consistency
- syncing backup controllers repeatedly, possible inconsistency/loss of state using long intervals

Network Virtualization

Network Virtualization

- abstraction of the physical network
- support for multiple logical networks running on a common shared physical substrate
- container of network services
- aspects of networks that can be virtualized:
 - nodes: virtual machines (fully-fledged VMs - Xen, KVM, VMware, VirtualBox; LXC)
 - links: tunnels (VLAN, VXLAN, GRE, Open vSwitch etc.)
 - storage:
 - block (Fibre Channel, iSCSI, SAS)
 - file (NFS, CIFS/SMB, AFS)



Network Virtualization - Motivation

- 2000s: lots of work on overlay networks on top of IP (here to stay)
- realization that one-size-fits all architectures are difficult
- Why not allow easier evolution?
- Promises:
 - rapid innovation: services delivered at software speeds
 - new forms of network control
 - vendor independence
 - simplified programming and operations
- Distinction:
 - SDN does not inherently abstract the details of the physical network
 - it does separate control and data plane
 - network virtualization provides the separation of logical and physical networks
 - SDN can be a useful tool for implementing virtual networks

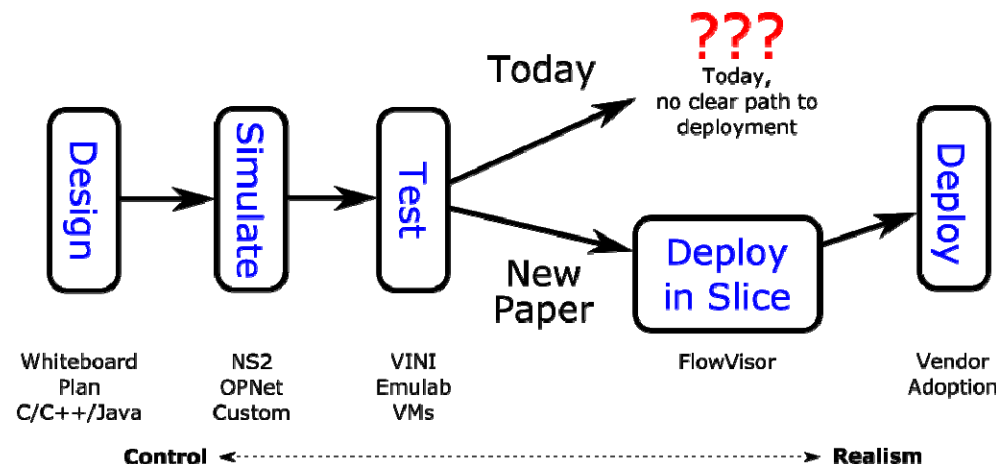
Network Virtualization - Design goals

1. Flexibility: topologies, routing and forwarding architecture; independent configuration
2. Manageability: separate policy and mechanisms
3. Scalability: maximize number of co-existing virtual networks
4. Security and Isolation: isolate both the logical networks and resources
5. Programmability: customizable/reprogrammable router, etc.
6. Heterogeneity: technology-agnostic, support for different technologies

Network Virtualization - Examples and Applications

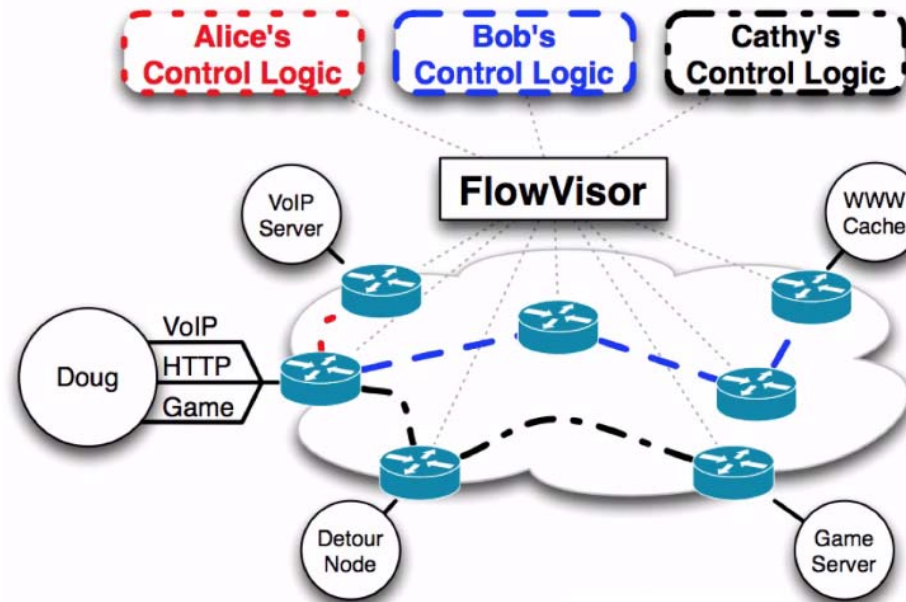
1) Experimentation on production networks

- How to test and deploy a "paper" design? → Goal: Realism
- Ideally: Deploy in parallel on production network



- FlowVisor: experimental traffic runs in parallel to production network
- different flows of user ("Doug") can be controlled by different controller designs (i.e. experimental protocols/architectures)
- a subset of flows (uncritical) can be controlled by new designs
- → virtualization on flow space (5 tuple: IP addresses, ports, protocol)

Network Virtualization - Examples and Applications



2) Dynamic scaling of resources

- own datacenter with limited resources
- fluctuation on demand, disaster, DDoS → additional resources required
- dynamically provision additional resources on demand as a leased service (dynamic scaling)
- Example: Amazon Virtual Private Cloud, EC2
 - allows customer to define own network, address space, etc.
 - extend existing enterprise data center

Network Virtualization - Examples and Applications

3) Network function virtualization


- unification of middlebox function
- today: purchase a variety of middleboxes separately: firewall, load-balancer, DPI, IDS/IPS
- instead:
 - distributed compute pool
 - dynamically install this functions as software (potentially on VMs)
 - interlinking via Virtual Networks

4) Rapid deployment and development of new network services

Overview of Control Plane

Overview of Control Plane - OpenFlow Specification

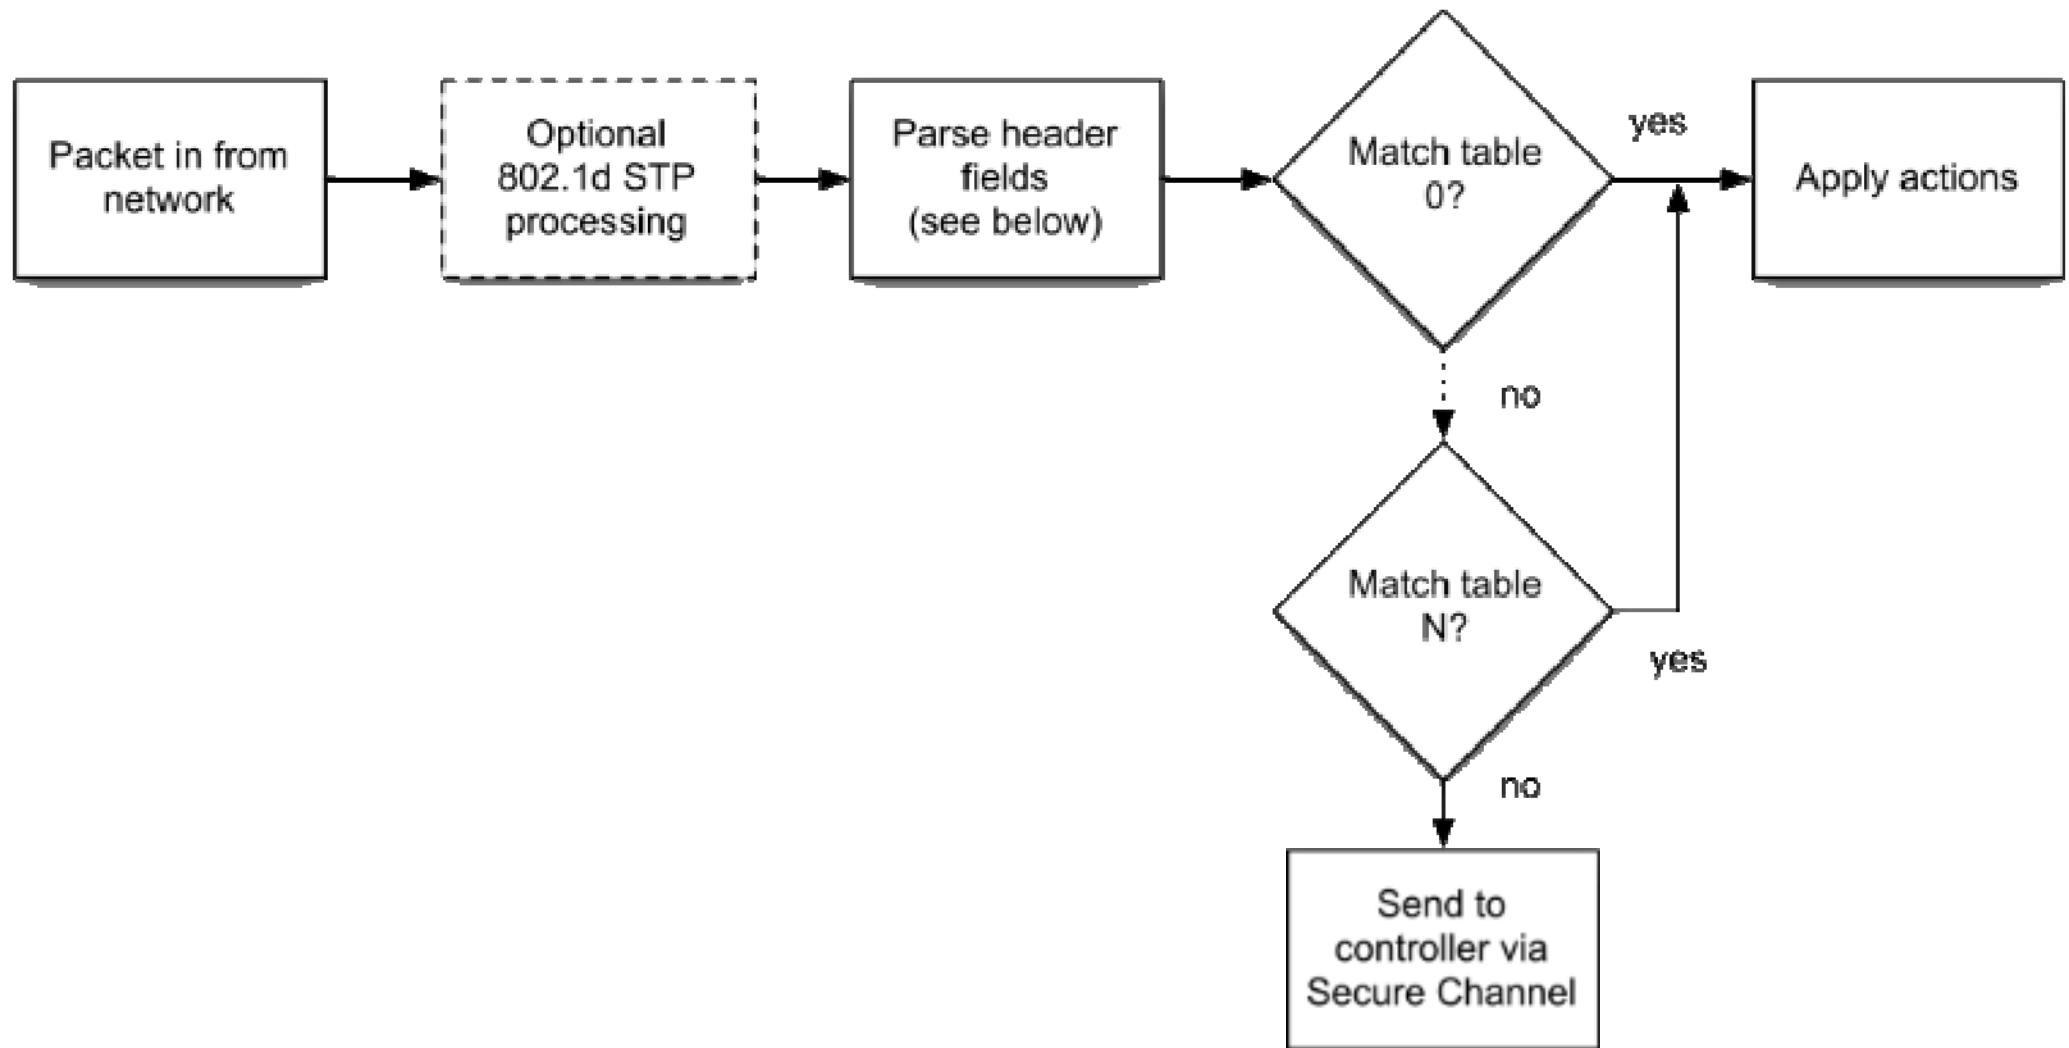
- OpenFlow controller communicates with switch over a secure channel
- OpenFlow protocol defines components of the switch, message format, types of actions the flow table should be able to perform
- purpose of control channel: update flow table entries in switch
- controller centric: logic executed at controller
- switch only forwards traffic based on flow table (data plane)



Overview of Control Plane - OpenFlow Switch Components

- Flow table: Performs packet lookup
 - matching the headers of incoming packets to flow table
 - performing action based on match being found (forward, drop, modify, enqueue)
 - no match → traffic is sent to controller
- Switch matches subset of packet header fields (12-tuple, wildcards):
 - Switch port
 - MAC src/dst
 - Ethernet type
 - VLAN ID/priority
 - IP src/dst
 - IP protocol
 - IP ToS bits
 - TCP/UDP sport/dport
- Secure channel: Communication to external controller
 - listens on control port 6634
 - inspecting flow table entries, modify flows, etc. (e.g. dpctl userspace program)

Overview of Control Plane - Matching




Overview of Control Plane - Mandatory Actions

- must be supported by OpenFlow v1.0 compliant switches
- Forward:
 - ALL: send out on all interfaces, except the incoming interface
 - CONTROLLER: encapsulate and send to controller
 - LOCAL: send to switch's local networking stack
 - TABLE: perform actions in flow table (packet-out messages)
 - IN PORT: sent packet out to the port, that packet was incoming from
 - optional: normal forwarding (based on routing table), spanning tree
- Drop:
 - flow-entry with no specified action → drop all matching packets

Overview of Control Plane - Optional Actions

- may be supported by OpenFlow v1.0 compliant switches
- Modify-Field : modify packet header values of the packet, e.g.
 - VLAN ID → redirect to logically separate networks
 - destination IP address → load balancing
- Enqueue: send packet through a queue attached to a port
 - apply QoS
 - traffic shaping

Overview of Control Plane - OpenFlow v1.3 Enhancements



- Action set: perform a set of actions on each matching packet
- Group: a list of action sets, allows switch to refer to a common set of actions performed on multiple sets of matching flows
- each table can update fields, modify action set → "Execute Action Set" performed before packet leaves
- recent version: v1.4 (Oct. 2013), OpenFlow is evolving

Overview of Control Plane - Action Group Options

- execute all action sets in a group
- e.g. implementing multicast: one packet is cloned for each action set in the group

Indirect group:

- one action set in a group is executed: performing same set of operations on multiple flow entries

Example actions:

- TTL: decrement, copy
- MPLS: apply (push) MPLS tags to a packet
- QoS: apply QoS actions (e.g. set_queue) to a packet
- metering and traffic monitoring

Overview of Control Plane - Other SDN Architectures

- Juniper's Contrail Controller
 - Linux-based
 - XMPP as control plane
 - L2 and L3 virtual networks
 - Contribution to OpenDaylight (FOSS implementation of various SDN control architectures)
- Cisco's Open Network Environment
 - centralized software controller
 - programmable data plane
 - ability to provide virtual overlays
- OpenFlow by far the most common SDN control architecture

Overview of Control Plane - OpenFlow Controllers

- only cover "Southbound interface": control channel between SDN controller and switches
- "Northbound interface": higher level of abstraction, policy layers on top of lower level SDN channels

1) NOX/POX

- first-gen OpenFlow controller
- FOSS, stable, widely used
- high-performance, clean codebase, well maintained and supported
- users implement control logic C++
- supports OpenFlow v1.0 (fork CPqD up to v1.3)
- model: controller registers for events → programmer writes event handlers
- low-level facilities and semantics of OpenFlow (low level of abstractions)
- POX (Python equivalent): easy to read/write code, worse performance

Overview of Control Plane - OpenFlow Controllers

2) Ryu


- implemented in Python → low performance
- supports OpenFlow up to v1.3
- works with OpenStack
- aims to be "Operating System" for SDN

3) Floodlight

- OpenFlow v1.0
- implemented in Java
- good documentation
- integration with REST API (RPC)
- production-level performance
- OpenStack integration
- disadvantage: steep learning curve

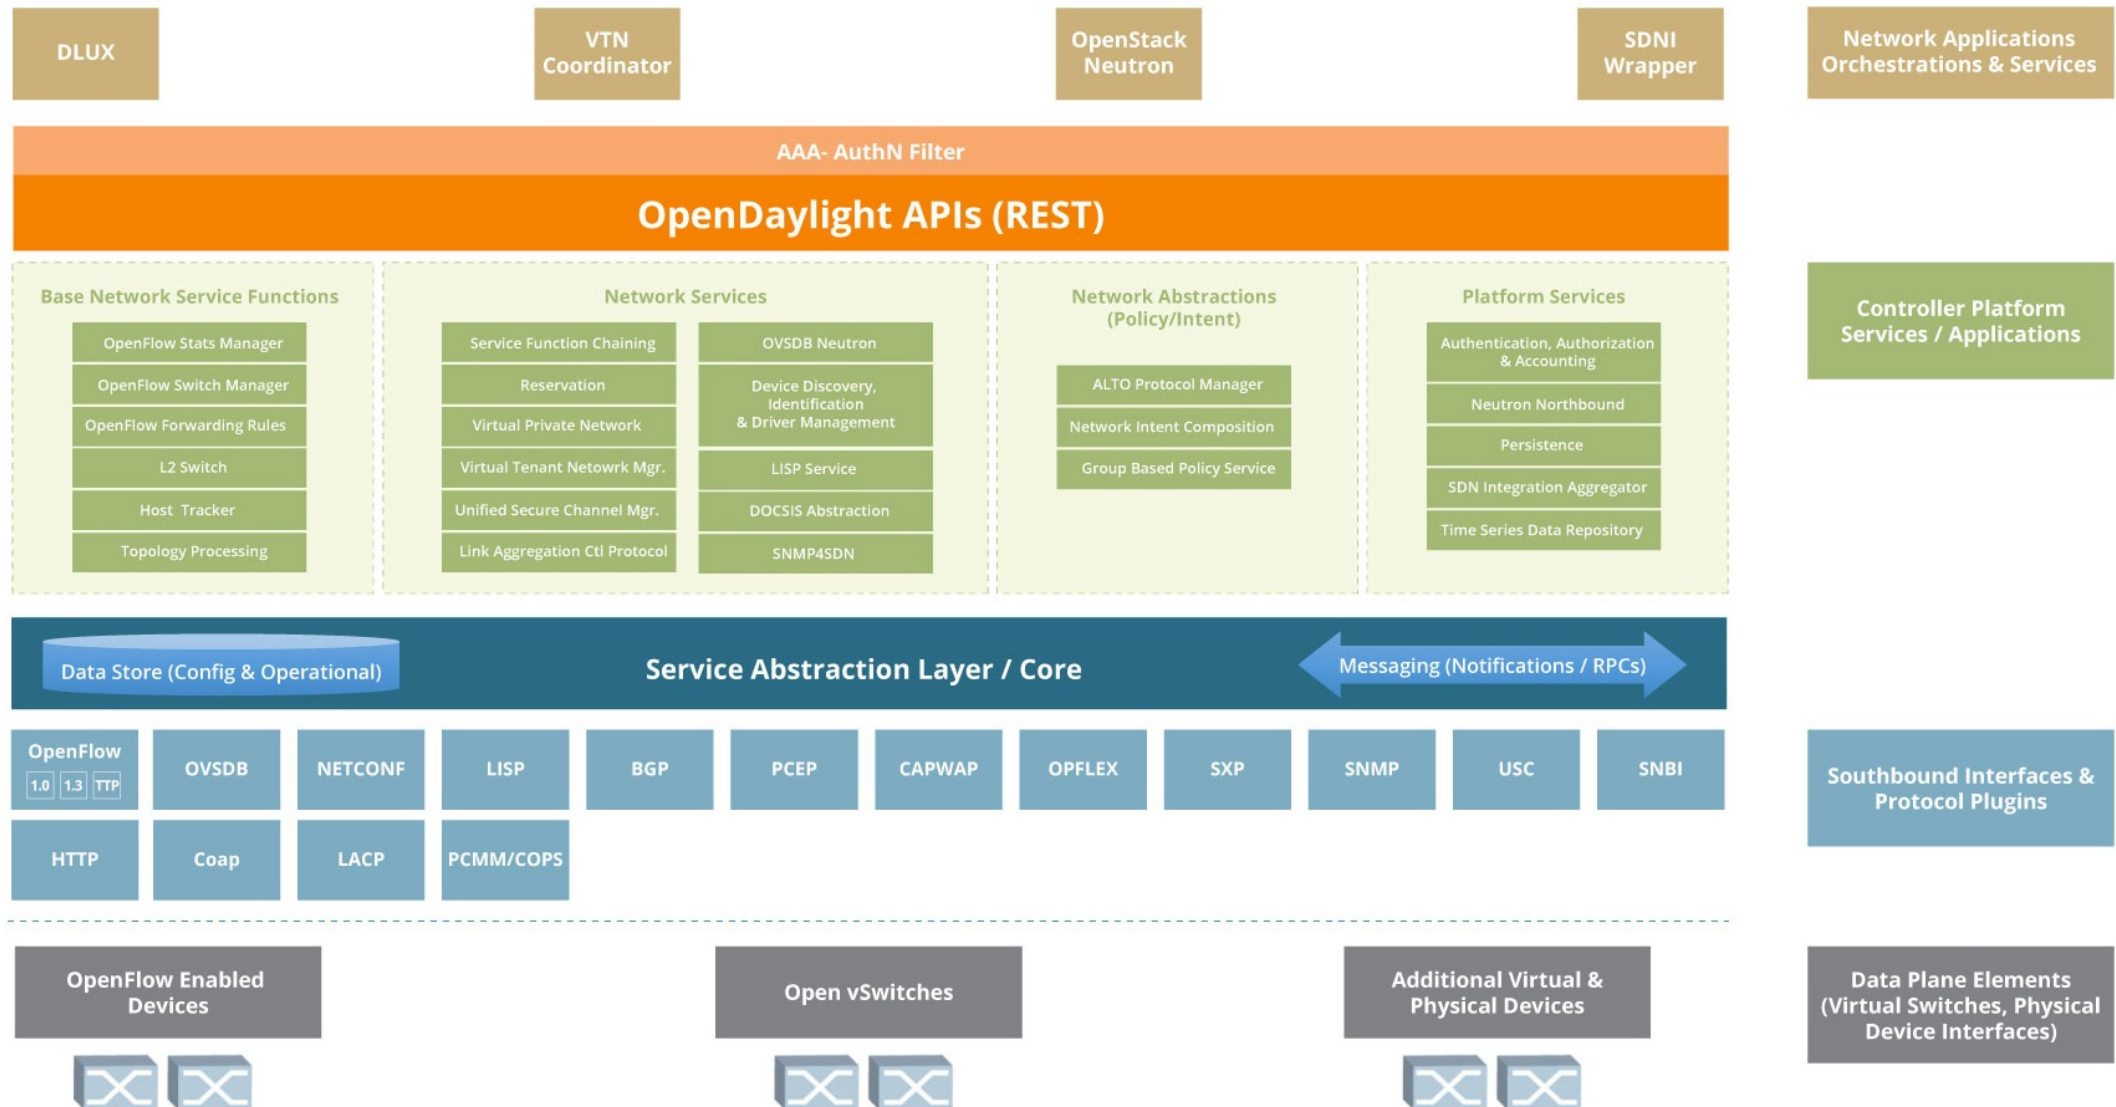
Overview of Control Plane - OpenFlow Controllers

3) OpenDaylight



- implemented in Java
- robust, extensible FOSS codebase
- heavy industry involvement and backing (8 platinum, 1 gold, 42 silver)
- common abstraction for northbound capabilities
- focus: open framework for building upon SDN/NFV innovations
- not limited to OpenFlow
- advantage: industry acceptance, integration with OpenStack, cloud applications, etc.
- disadvantage: complex → steep learning curve, rather poor documentation

Overview of Control Plane - OpenDaylight Architecture



Overview of Control Plane - Controller Comparison

	NOX	POX	Ryu	Floodlight	ODL
Language	C++	Python	Python	Java	Java
Performance	Fast	Slow	Slow	Fast	Fast
Distributed	No	No	Yes	Yes	Yes
OpenFlow	1.0	1.0	1.0, 1.1, 1.3, 1.4	1.0	1.0, 1.3
Multi-tenant clouds	No	No	Yes	Yes	Yes
Learning curve	Moderate	Easy	Moderate	Steep	Steep

Control Plane - Motivation for "Northbound" Interface

OpenFlow programming is not easy:

- low level of abstraction (channel to control flow table entries)
- difficult to implement higher level task (e.g. security policies, load balancing) using OpenFlow
- difficult to perform multiple independent tasks (e.g. interfering/conflicting routing and access control)
- only unhandled packets inspected by controller → incomplete view
- race conditions, if flow tables rules are not installed properly, e.g. "routing" loops

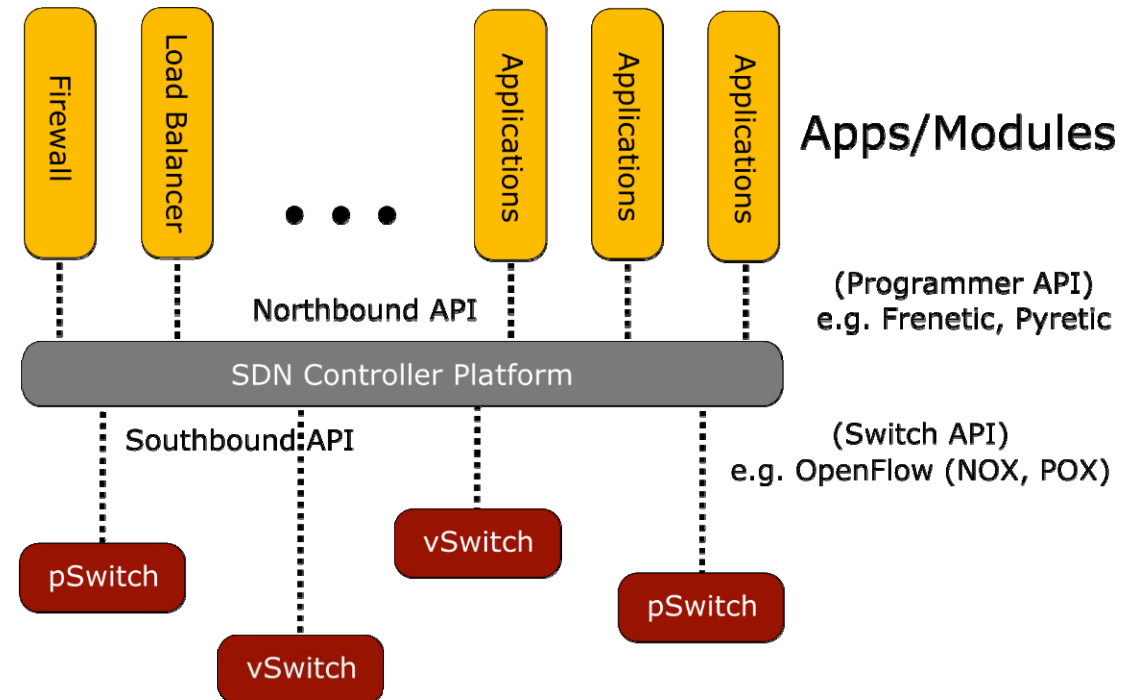
→ "Northbound" interface:

- programming interface allows applications and orchestration systems to program the network
- higher level of abstraction
- → policies are being "compiled" into OpenFlow rules

Control Plane - Motivation for "Northbound" Interface

Use cases:

- path computation
- routing
- recover from failures
- security policies



Users of the "Northbound" interface:

- sophisticated network operators
- service providers: value-added services
- vendors: create services on top of a switch/controller
- researches
- anyone, who wants to develop capabilities on top of OpenFlow

Control Plane - Motivation for "Northbound" Interface

Benefits:

- vendor independence
- quickly modify/customize control through popular programming languages

Examples:

- large virtual switch
- security applications
- resource management and control (traffic engineering/load balancing)
- middlebox integration

Network Functions Virtualization

Network Functions Virtualization

- place arbitrary functions in VMs and distribute them across the network
- status quo: middlebox functions (firewall, load-balancer, DPI, IDS/IPS) placed in separate, monolithic middleboxes
- NFV: functions distributed in VMs/containers across the network
- decoupling functions from hardware → increased flexibility how packet processing is performed

Benefits:

- reduced CAPEX/OPEX, reduced time to market
- elastic scaling
- vendor agnostic

New use cases:

- virtualized services for enterprises, virtual CDNs
- virtualized mobile core network (decreased latency)
- integration of production and testing

Network Functions Virtualization

- fine-grained functional elements, instead of monolithic middleboxes, e.g.
- WAN optimizer = Caching + Deduplication + Compression + Encryption + FEC + Rate Limiter
- Application firewall = IP-Defragmenter + Application Detection Engine + Logger + Blocker
- → placing individual functional elements in virtual containers and chain them according to middlebox functionality, reuse of functional elements
- Problem: Orchestration and customizability
 - enable network operators to implement modular network functions without worrying about Placement and Steering
 - add custom middlebox functions inside network data plane
- Future work:
 - (better) algorithms for placement and steering
 - high-speed data plane implementation: throughput, instantiation, migration
 - better ways to resolve policy conflicts

NFV - Difficulties in Placement and Steering

Placement: Where to place functions in the network? Goals:

- minimize the number of locations to place the elements
- minimize the bandwidth utilization for implementing the policy
- minimize the latency of flows for implementing the policy

Steering: How to route traffic through these functions? Goals:

- not only shortest path
- consistent ordering
- chaining in a certain order
- dynamic chaining: enable path changes for packet flow based on certain conditions (e.g. stateful firewalls)
- load awareness

NFV - Difficulties in Placement and Steering

Mapping traffic flows and demands to:

- available network resources (i.e. paths)
 - available processing capacity (i.e. middleboxes)
-
- need for a unified abstraction for control, data and storage
 - continuously monitor link and machine loads
 - create new element instances in case of overload
 - steer some traffic through new instances
 - reclaim unused element instances